

Homework 1: Writing Your Own Shell

Assigned: January 21, 2015

Due: February 5, 2015 (by class time)

1. Assignment

Your boss requires a lean-and-mean UNIX shell, and none of the existing shells (sh, bash, ksh, csh, tcsh, zsh, ash, or even the infamous “adventure shell”) seem to meet her requirements. So, you are stuck writing a shell almost from scratch! It does not need to support job management, environment variables, or most other standard shell features. However, it must be able to run programs and support file redirection and pipes.

2. Shell Specification

This section outlines your boss’s requirements for the shell that you are writing. It is fine to discuss the specification and its interpretation with other students, but do not share code. You should ask for clarification if you think you have found an error or ambiguity in the spec. Don’t be put off by the length of the specification: it is simply formalizing a simple version of the UNIX shell syntax that you have always used.

3. The Shell Language

Lexical structure:

- The input to the shell is a sequence of **lines**. The shell must correctly handle lines of up to 100 characters. If a line containing more than 100 characters is submitted to the shell, it should print some kind of error message and then continue processing input at the start of the next line.
- Each line consists of **tokens**. Tokens are separated by one or more spaces. A line may contain as many tokens as can fit into 100 characters.
- There are two kinds of tokens: **operators** and **words**. The only operators are `<`, `>`, and `|` (pipe).
- Words consist of the characters A–Z, a–z, 0–9, dash, dot, forward slash, and underscore. If a word in a line of input to the shell contains any character not in this set, then the shell should print an error message and then continue processing input at the start of the next line.
- The only legal input to the shell, other than lines consisting of valid tokens, is the end-of-file.

Parsing the shell language:

- Lines of input are divided into token groups. Each token group will result in the shell forking a new process and then executing a program.
- Every token group must begin with a word that is called the command. The words immediately following a command are called arguments and each argument belongs to the command it most closely follows. The order of arguments matters, and they are numbered from left to right, starting at one.
- It is permissible for the arguments in a token group to be followed by one or more file redirections. A file redirection consists of one of the operators `<` or `>` followed by a single word called the filespec. A file redirection containing the `<` operator is an **input** file redirection and a file redirection containing the `>` operator is an **output** file redirection.
- Token groups are separated by pipe operators. In other words, each valid line of shell input must begin and end with a valid token group, and the only place pipe operators are allowed is in between token groups.
- Valid token groups may be preceded by a pipe operator or they may contain an input file redirection, or neither, but not both. Similarly, valid token groups may be followed by a pipe operator or contain an output file redirection, or neither, but not both.
- Lines of shell input that violate any of the parsing rules should cause the shell to print an error message and then move on to the next line of input.

Here are a few examples of shell language:

- `/usr/bin/emacs`
is a legal line of input. It contains a single token group containing a single token, which is a command.
- `/usr/bin/emacs |`
is not a legal line of input because the pipe character is not valid as part of a **word** (remember that tokens are always separated by one or more spaces).
- `/usr/bin/emacs |`
is not a legal line of input because the pipe operator has to separate valid token groups. It is not legal for a pipe operator to be at the end of a line of input.
- `ls -l > foo`
is a legal line of input containing a single token group. In order, the tokens are command, argument, operator, filespec.
- `ls -l > | foo2`

is not a legal line of input because the output redirection operator is not followed by a word

- `ls -l > foo1 | foo2`

is not a legal line of input because the token group is both followed by a pipe and contains an output redirection.

- `> foo`

is not a legal line of input because it does not begin with a word.

- `ls > foo%`

is not a legal line of input because it contains an illegal character.

- `ls | grep -i cs5460 | sort | uniq | cut -c 5`

is a legal line of input. It contains 5 token groups. In order, the tokens in this line are: command, operator, command, argument, argument, operator, command, operator, command, operator, command, argument, argument.

4. Interpreting the Shell Language

Only legal lines of input (as defined in the previous section) should be interpreted — when the shell encounters an illegal line it prints an error message and then continues to the next line of input. (Again, your co-worker's parser handles the error messages.)

- Every **command** except the special command `exit` is to be interpreted as a UNIX executable to be `exec'd`.
- When the shell encounters either the command `exit` or the end-of-file, it terminates with a return code of zero.
- When a command begins with the character `'/'` it is an absolute pathname and should be `exec'd` as-is. All other commands are **relative** and should be assumed to be in the current working directory (CWD), and therefore the shell must prepend the name of the CWD to the command before `execing` it.
- The arguments to a command should be passed to the `exec` call in `argv`; `argv[0]` should always be the same as the string that is being `exec'd`, with actual arguments passed in slots one and higher.
- The `>` operator indicates that `STDOUT` of the associated command should be redirected to the UNIX file named by the filespec belonging to the operator. This file should be created if it does not exist, and the shell should report an error if the file cannot be created. Similarly, the `<` operator indicates that `STDIN` of the associated command should be redirected from the UNIX file named by the filespec. The shell should report an error if this file cannot be opened for reading.

- The pipe (`|`) operator indicates that `STDOUT` of the preceding command should be redirected to `STDIN` of the following command.
- After interpreting a command, the shell should wait for all forked subprocesses to terminate before parsing the next line of input. Also, after all forked subprocesses have terminated, the shell should report their return values (e.g., one per line in the order that the commands were specified in the line of input).

Here is the interpretation of the legal example commands from the previous section. Assume that the CWD is `/home/regehr`.

- `/usr/bin/emacs`

The shell forks a new process and in it execs `/usr/bin/emacs`; the main shell process waits for `emacs` to exit before reading another line of input.

- `ls -l > foo`

The shell opens the file `foo` for writing, forks a new process, redirects `STDOUT` of the new process to `foo`, and then execs `/home/regehr/ls` with `-l` as argument one.

- `ls | grep -i cs4414 | sort | uniq | cut -c 5`

The shell sets up the pipes, forks 5 subprocesses, and then each subprocess execs the command for one of the 5 token groups.

5. Getting Started

Be sure you understand this assignment before starting to write code. Here is a rough outline of steps you might take in solving it.

1. *Command-Line Parsing*

Parse each input line into an array of strings.

2. *Interpreting Shell Commands*

Modify the shell command loop to properly interpret lines of input. You will need to use the `fork()` and `execve()` functions to do this. The shell process should wait for its children to complete by calling `waitpid()`. The shell should also check and report the exit code returned by the program. Note: the exit code is placed into the lower 8-bits of the status code set by `waitpid()`. The full version of your code will look roughly like this:

```
while (1) {
    read a line of input
    parse the line
    for each command in the line {
        pid = fork();
```

```

    if (pid == 0) {
        do redirection stuff
        execve ( command, args , ...);
        oops, why did exec fail?
    } else {
        store pid somewhere
    }
}
for each command in the line {
    waitpid (stored pid, &status);
    check return code placed in status;
}
}

```

But remember to practice **incremental development**: get something working, add a small feature, test your code, and then move on to the next feature. For example, you should start out supporting shell input that contains a single command and no redirection. Then, add one of these features and then the other.

Your shell should return results similar to those returned by well-known shells like `tcsh` and `bash`.

To support relative pathnames, use the `getcwd()` system call to find the name of the directory in which your shell is running. You can make this call once at startup time since your shell does not need to support changing directories.

To support **I/O redirection**, modify the child process created by `fork()` by adding some code to open the input and output files specified on the command line. This should be done using the `open()` system call. Next, use the `dup2()` system call to replace the standard input or standard output streams with the appropriate file that was just opened. Finally, call `exec()` to run the program.

Pipes are a little trickier: you should use the `pipe()` system call to create a pair of pipe file descriptors **before** calling `fork()`. After the fork both processes will have access to both sides of the pipe. The reading process should close the write file descriptor, and the writing process should close the read file descriptor. At this point each process uses `dup2()` to copy the remaining pipe descriptor over `STDIN` or `STDOUT` as appropriate.

When you are done, bask in the glory of a working shell! You should now have a good operational understanding of the user-mode side of some of the most important UNIX system calls.

3. *Other Odds and Ends*

Finding Commands: Since your shell is not implementing search paths, typing absolute paths to programs may become tiresome. One easy way around this is to create some symbolic links from your working directory to commonly used programs. For example:

```
ln -s /bin/ls ./ls
```

Code Size: John Regehr’s finished shell was under 400 lines of C code, including plenty of debugging code, blank lines, and comments. Your shell should not be a lot larger than this (in other words, if it gets a lot larger then you’re probably doing something wrong).

Getting Help: Since this is an upper division computer science course, you are expected to do your own research regarding the usage of various system calls, header files, and libraries. Information is readily available in the man pages, UNIX reference books, and on the web. For example, on any UNIX machine “man pipe” will give you information about the pipe() system call. Otherwise, do not hesitate to ask a question if you are unclear about how some part of the assignment is supposed to work.

6. Logistics and Grading

Remember, this assignment is **not** a group assignment — every individual in the class should do it themselves (and hand it in themselves)! Also, remember to start early, it’s easy to get hung up on problems that would seem easy if you had more time to read documentation and think about them.

Although you may do your work on any UNIX machine, your shell will be graded on our Linux VM, so you must make sure that it compiles and runs properly there.

You will hand in via collab a tar file into collab with all of your sources and headers, a *makefile*, and a PDF of your write-up. Your write-up will include the source files as well. We will un-tar your tarfile into a directory, *cd* into the directory, *rm msh*, and run *make*. Your make file should do all necessary compilation to create your shell. Your shell will be called *msh*. In other words, after make is called there will be an executable called *msh*. We will then proceed to test your shell.

You will also hand in a printed copy of your write-up in class.